

AgentShadow

Discover, score & govern every AI agent in your organization

The problem

AI agents are proliferating across codebases and SaaS stacks - LangChain, CrewAI, AutoGPT, AutoGen, OpenAI Assistants. Most are invisible to security teams: no inventory, no risk rating, no governance. A single shell-capable, unowned, high-autonomy agent is a material breach risk.

What AgentShadow does

- **Discover** agents via code scanning *and* runtime/SaaS connectors.
- **Inventory** them in a single Discovered Agents view (owner, framework, tools, autonomy).
- **Score** each agent 0-100 with a deterministic, auditable risk engine.
- **Govern** with YAML policies: allow / warn / deny, with full decision trails.
- **Report** branded PDF assessments and feed the cross-tool correlation graph.

What we detect

Frameworks	LangChain/LangGraph, CrewAI, AutoGPT, AutoGen, OpenAI Assistants, LlamaIndex, Semantic Kernel
Risk signals	Shell/code execution, unbounded autonomy, missing approval gates, hardcoded secrets, exfiltration tools
Governance	Score thresholds, capability gates, ownership/accountability requirements

Built on a proven platform

AgentShadow reuses battle-tested engines from the Valo security ecosystem - LLMSHadow framework detection, the Valo deterministic scoring and governance policy engines, Valo ReportLab reporting, and SaaSShadow inventory UI - registering as a first-class source in the shared correlation graph. New detection content, not a new architecture.